

## Day 3 : Statistical analysis using R

### Session 9: Regression

#### 9.1. Fitting regression models with `lm()`

Symbols commonly used in R formulas

Symbol	Usage
~	Separates response variables on the left from the explanatory variables on the right. For example, a prediction of $y$ from $x$ , $z$ , and $w$ would be coded $y \sim x + z + w$ .
+	Separates predictor variables
:	Denotes an interaction between predictor variables. A prediction of $y$ from $x$ , $z$ , and the interaction between $x$ and $z$ would be coded $y \sim x + z + x:z$ .
*	A shortcut for denoting all possible interactions. The code $y \sim x * z * w$ expands to $y \sim x + z + w + x:z + x:w + z:w + x:z:w$ .
^	Denotes interactions up to a specified degree. The code $y \sim (x + z + w)^2$ expands to $y \sim x + z + w + x:z + x:w + z:w$ .
.	A placeholder for all other variables in the data frame except the dependent variable. For example, if a data frame contained the variables $x$ , $y$ , $z$ , and $w$ , then the code $y \sim .$ would expand to $y \sim x + z + w$ .
-	A minus sign removes a variable from the equation. For example, $y \sim (x + z + w)^2 - x:w$ expands to $y \sim x + z + w + x:z + z:w$ .
-1	Suppresses the intercept. For example, the formula $y \sim x - 1$ fits a regression of $y$ on $x$ and forces the line through the origin at $x=0$ .
I ()	Elements within the parentheses are interpreted arithmetically. For example, $y \sim x + (z + w)^2$ expands to $y \sim x + z + w + z:w$ . In contrast, the code $y \sim x + I((z + w)^2)$ expands to $y \sim x + h$ , where $h$ is a new variable created by squaring the sum of $z$ and $w$ .
function	Mathematical functions can be used in formulas. For example, $\log(y) \sim x + z + w$ predicts $\log(y)$ from $x$ , $z$ , and $w$ .

#### E034-simple\_regression.R

```
# Set seed for reproducibility
set.seed(12345)

# Generate 1000 observations
n <- 1000

# Generate study_hours as uniform random numbers between 0 and 10
study_hours <- round(runif(n, min = 0, max = 10))

# Generate score as a linear function of study_hours with noise
score <- 50 + 5 * study_hours + rnorm(n, mean = 0, sd = 5)

# Combine into a data frame
df <- data.frame(study_hours, score)

# Perform linear regression
model <- lm(score ~ study_hours, data = df)

# Summarize the regression results
summary(model)
```

① simple regression with generated data

```
Call:
lm(formula = score ~ study_hours, data = df)


Residuals:
    Min       1Q   Median       3Q      Max
-16.612  -3.334  -0.018   3.509  16.737

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  49.94446    0.32742   152.54  <2e-16 ***
study_hours   4.99446    0.05571    89.64  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.035 on 998 degrees of freedom
Multiple R-squared:  0.8895,    Adjusted R-squared:  0.8894
F-statistic: 8036 on 1 and 998 DF, p-value: < 2.2e-16
```

② visualizing simple linear regression

```
# visualizing simple regression
library(ggplot2)
ggplot(data = df, aes(x=study_hours, y=score)) +
  geom_point() +
  geom_smooth(method = 'lm') +
  theme_bw()
```



## 9.2. Multiple regression

### E035-multiple\_regression.R

```
# Set seed for reproducibility
set.seed(12345)

# Generate 200 observations
n <- 200

# Generate age variable (cycles from 18 to 69)
age <- (1:n %>% 52) + 18

# Generate educ_year variable (cycles from 0 to 17)
educ_year <- (1:n %>% 18)

# Generate income variable with a linear relationship to age and educ_year,
plus noise
income <- 20000 + 800 * age + 3000 * educ_year + rnorm(n, mean = 0, sd =
2000)
```

```
# Combine into a data frame
```

```
df <- data.frame(age, educ_year, income)
```

① Data generation

```
# Regression with omitted variable
```

```
model_omitted <- lm(income ~ age, data = df)
```

```
summary(model_omitted)
```

```
Call:
```

```
lm(formula = income ~ age, data = df)
```

```
Residuals:
```

```
      Min       1Q   Median       3Q      Max
-29181.6 -13567.7   363.5  14563.5  27208.4
```

```
Coefficients:
```

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  44039.53    3535.95   12.46  <2e-16 ***
age           835.97      78.13    10.70  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 16120 on 198 degrees of freedom
```

```
Multiple R-squared:  0.3663,    Adjusted R-squared:  0.3631
```


```
F-statistic: 114.5 on 1 and 198 DF,  p-value: < 2.2e-16
```

② Regression with omitted variable

```
# Residual diagnostics for omitted variable model
```

```
plot(model_omitted, which = 1) # Residual vs Fitted plot
```


③ Residual diagnostic (visual)



```
resid_omitted <- residuals(model_omitted)
```

```
hist(resid_omitted)
```

Histogram of resid\_omitted



```
shapiro.test(resid_omitted), # Shapiro-Wilk test for normality [H0: normally distributed]
```

④ residual normality test.

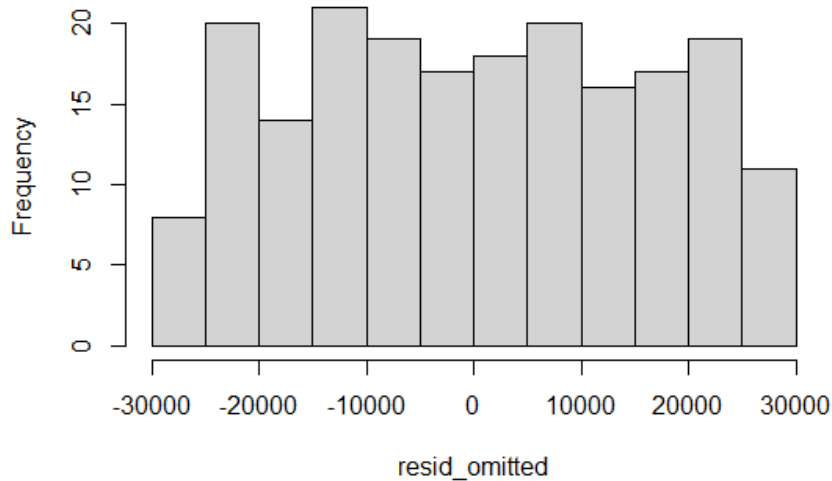
### Shapiro-Wilk normality test

```
data: resid_omitted  
W = 0.95517, p-value = 6.14e-06
```


⑤ multiple regression with correct specification.

```
#-----  
# Multiple regression with correct specification  
model_correct <- lm(income ~ age + educ_year, data = df)  
summary(model_correct)
```


Histogram of resid\_omitted



```
# Residual diagnostics for correctly specified model  
plot(model_correct, which = 1) # Residual vs Fitted plot
```



```
lm(income ~ age + educ_year)  
resid_correct <- residuals(model_correct)  
hist(resid_correct)
```



### 9.3. Polynomial regression

```
E036-polynomial_regression.R
library(ggplot2)

mtcars <- datasets::mtcars

#-----
#simple regression
#-----


fit <- lm(data = mtcars, formula = mpg ~ hp) # mpg: Miles/(US) gallon, hp:
Gross horsepower
summary(fit) #R-squared : 0.6024, Residual standard error: 3.863
Call:
lm(formula = mpg ~ hp, data = mtcars)

Residuals:
    Min       1Q   Median       3Q      Max
-5.7121 -2.1122 -0.8854  1.5819  8.2360

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 30.09886   1.63392  18.421 < 2e-16 ***
hp          -0.06823   0.01012  -6.742 1.79e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.863 on 30 degrees of freedom
Multiple R-squared:  0.6024,    Adjusted R-squared:  0.5892
F-statistic: 45.46 on 1 and 30 DF, p-value: 1.788e-07

ggplot(mtcars, aes(x = hp, y = mpg)) +
  geom_point() +
  stat_smooth(method = 'lm', formula = y ~ x, color = 'red', se = FALSE) +
  theme_bw()
```



②

```
#-----
#Polynomial regression
#-----
fit <- lm(data = mtcars, formula = mpg ~ hp + I(hp^2))
summary(fit) #R-squared : 0.7561, Residual standard error: 3.077
Call:
lm(formula = mpg ~ hp + I(hp^2), data = mtcars)


Residuals:
    Min       1Q   Median       3Q      Max
-4.5512 -1.6027 -0.6977  1.5509  8.7213

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  4.041e+01  2.741e+00  14.744 5.23e-15 ***
hp          -2.133e-01  3.488e-02  -6.115 1.16e-06 ***
I(hp^2)       4.208e-04  9.844e-05   4.275 0.000189 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.077 on 29 degrees of freedom
Multiple R-squared:  0.7561,    Adjusted R-squared:  0.7393
F-statistic: 44.95 on 2 and 29 DF,  p-value: 1.301e-09
```

upward/downward slopping?  
U-shaped or inverted U-shaped?

```
ggplot(mtcars, aes(x = hp, y = mpg)) +
  geom_point() +
  stat_smooth(method = 'lm', formula = y ~ x + I(x^2), color = 'red', se =
FALSE) +
  theme_bw()
```



## 9.4. Regression with interaction term

a. Interaction term can be defined as the multiplication of two variables.

### E037-regression\_with\_interaction.R

```
mtcars <- datasets::mtcars
```

②

```
#generating a new interaction term hp * wt
mtcars$hp_wt <- mtcars$hp * mtcars$wt

fit <- lm(mpg ~ hp + wt + hp_wt, data=mtcars)
summary(fit)

Call:
lm(formula = mpg ~ hp + wt + hp_wt, data = mtcars)

Residuals:
    Min       1Q   Median       3Q      Max
-3.0632 -1.6491 -0.7362  1.4211  4.5513

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  49.80842     3.60516   13.816 5.01e-14 ***
hp           -0.12010     0.02470    -4.863 4.04e-05 ***
wt           -8.21662     1.26971    -6.471 5.20e-07 ***
hp_wt         0.02785     0.00742     3.753 0.000811 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.153 on 28 degrees of freedom
Multiple R-squared:  0.8848,    Adjusted R-squared:  0.8724
F-statistic: 71.66 on 3 and 28 DF,  p-value: 2.981e-13
```

b. Interaction term is used when the effect of one variable depends on the value of another variable.

c. e.g. wage ~ experience \* age

#OR

```
fit <- lm(mpg ~ hp + wt + hp:wt, data=mtcars)
summary(fit)

Call:
lm(formula = mpg ~ hp + wt + hp:wt, data = mtcars)

Residuals:
    Min       1Q   Median       3Q      Max
-3.0632 -1.6491 -0.7362  1.4211  4.5513

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  49.80842     3.60516   13.816 5.01e-14 ***
hp           -0.12010     0.02470    -4.863 4.04e-05 ***
wt           -8.21662     1.26971    -6.471 5.20e-07 ***
hp:wt         0.02785     0.00742     3.753 0.000811 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.153 on 28 degrees of freedom
Multiple R-squared:  0.8848,    Adjusted R-squared:  0.8724
F-statistic: 71.66 on 3 and 28 DF,  p-value: 2.981e-13
```

③ Direct way of interaction term

```
##* -----
##* A significant coefficient of interaction term indicates that
##* the relationship between mpg and hp varies by wt. Similarly,
##* the relationship between mpg and wt varies by hp.
##* -----
```

④

```
# d(mpg)/d(hp) = - 0.12010 + 0.02785 * wt
wt = 1
print(- 0.12010 + 0.02785 * wt) #-0.09225

wt = 2
print(- 0.12010 + 0.02785 * wt) #-0.0644

wt = 3
print(- 0.12010 + 0.02785 * wt) #-0.03655

# d(mpg)/d(wt) = - 8.21662 + 0.02785 * hp
hp = 100
```

```
print(- 8.21662 + 0.02785 * hp) #-5.43162

hp = 150
print(- 8.21662 + 0.02785 * hp) #-4.03912

hp = 200
print(- 8.21662 + 0.02785 * hp) #-2.64662
```

## 9.5. Logarithmic regression

① OLS assumes linear relationship but in real word relationships are non-linear.  
- taking log help linearize relationship.

② four types of models

Model	Equation	Interpretation of $\beta_1$
Log-Log	$\log(y) = \beta_0 + \beta_1 \log(x)$	Elasticity: 1% change in $x$ leads to $\beta_1$ % change in $y$ .
Log-Linear	$\log(y) = \beta_0 + \beta_1 x$	Semi-elasticity: 1-unit change in $x$ leads to $(\exp(\beta_1) - 1) \times 100\%$ change in $y$ .
Linear-Log	$y = \beta_0 + \beta_1 \log(x)$	1% change in $x$ leads to $\beta_1/100$ unit change in $y$ .
Linear-Linear	$y = \beta_0 + \beta_1 x$	1-unit change in $x$ leads to $\beta_1$ unit change in $y$ .

### E038-logarithmic\_regression.R

```
#-----
# Log-Log Regression
#-----
# Load data
mtcars <- datasets::mtcars

② # Log-log regression
model_loglog <- lm(log(mpg) ~ log(displacement), data = mtcars)
summary(model_loglog)
Call:
lm(formula = log(mpg) ~ log(displacement), data = mtcars)

Residuals:
    Min       1Q   Median       3Q      Max
-0.22758 -0.08874 -0.00791  0.07970  0.32143

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  5.38097    0.20803   25.87 < 2e-16 ***
log(displacement) -0.45857    0.03913  -11.72 1.01e-12 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1282 on 30 degrees of freedom
Multiple R-squared:  0.8207,    Adjusted R-squared:  0.8148
F-statistic: 137.3 on 1 and 30 DF,  p-value: 1.006e-12

# A 1% increase in Displacement (cu.in.) reduces Miles/(US) gallon by ~0.46%.

#-----
④ # Log-Linear Regression
#-----
# Log-linear regression
model_loglin <- lm(log(mpg) ~ hp, data = mtcars)
summary(model_loglin)
```



```

Call:
lm(formula = log(mpg) ~ hp, data = mtcars)

Residuals:
    Min       1Q   Median       3Q      Max
-0.41577 -0.06583 -0.01737  0.09827  0.39621

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  3.4604669   0.0785838   44.035 < 2e-16 ***
hp          -0.0034287   0.0004867   -7.045 7.85e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1858 on 30 degrees of freedom
Multiple R-squared:  0.6233,    Adjusted R-squared:  0.6107
F-statistic: 49.63 on 1 and 30 DF,  p-value: 7.853e-08

# A 1-unit increase in horsepower reduces MPG by ~0.34% (exp(-0.0034287) - 1
≈ -0.003422829).

#-----
# Linear-Log Regression
#-----
# Load data
trees <- datasets::trees

# Linear-log regression
model_linlog <- lm(Volume ~ log(Girth), data = trees)
summary(model_linlog)
Call:
lm(formula = Volume ~ log(Girth), data = trees)

Residuals:
    Min       1Q   Median       3Q      Max
-9.7246 -3.5312 -0.9174  3.2154 15.8780

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -138.973     11.439  -12.15 6.71e-13 ***
log(Girth)   66.141       4.455   14.85 4.38e-15 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.701 on 29 degrees of freedom
Multiple R-squared:  0.8837,    Adjusted R-squared:  0.8797
F-statistic: 220.4 on 1 and 29 DF,  p-value: 4.381e-15

# A 1% increase in girth increases volume by ~ 0.66 units (66.141 / 100).

```

## Session 10: Logistic regression

### 11.1. Logistic regression

① When the dependent variable is a binary variable e.g. employed.

Logistic regression is useful when you're predicting a binary outcome from a set of continuous and/or categorical predictor variables.

#### E039-Logistic\_regression.R

```

# Load necessary libraries
library(haven) # For reading SPSS files
library(dplyr) # For data manipulation
library(margins) # For calculating marginal effects

# Import SPSS file from the URL ② loading data
data <- read_spss('data/010-hh.sav')

# Dropping missing values in HHSEX
data <- data %>% filter(!is.na(HHSEX))

# Creating new variables
data <- data %>%
  mutate(
    hh_size = HH48, # HH member size variable
    urb_rur = factor(HH6), # 1=Urban 2=Rural
    province = factor(HH7), # Province number

```

```

hhsex = factor(HHSEX) # 1=Male 2=Female
)

#setting 1=Urban as reference/base
data$urb_rur <- relevel(data$urb_rur, ref = '1')

#setting 2=Female as reference/base
data$hhsex <- relevel(data$hhsex, ref = '2')

#setting province 3 as base category/reference level
data$province <- relevel(data$province, ref = '3')

#-----
# Running logistic regression
#-----
logit_model <- glm(hhsex ~ hh_size + urb_rur + province,
                  data = data, family = binomial(link = "logit"))
summary(logit_model)
Call:
glm(formula = hhsex ~ hh_size + urb_rur + province, family = binomial(link = "logit"),
    data = data)

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -0.39332    0.06488  -6.062 1.35e-09 ***
hh_size      0.33308    0.01274  26.145 < 2e-16 ***
urb_rur2     0.16929    0.04273   3.962 7.44e-05 ***
province1    0.24989    0.07272   3.436 0.000590 ***
province2    0.46344    0.07949   5.830 5.53e-09 ***
province4   -0.47811    0.06912  -6.917 4.61e-12 ***
province5   -0.28721    0.06988  -4.110 3.95e-05 ***
province6   -0.22815    0.07692  -2.966 0.003017 **
province7   -0.25820    0.07476  -3.454 0.000553 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 14830  on 12654  degrees of freedom
Residual deviance: 13722  on 12646  degrees of freedom
AIC: 13740

Number of Fisher Scoring iterations: 4

# Calculating marginal effects for logistic regression
logit_margins <- margins(logit_model)
summary(logit_margins)
  factor      AME      SE      z      p    lower    upper
hh_size  0.0606  0.0021 28.5476 0.0000  0.0564  0.0647
province1 0.0424  0.0122  3.4806 0.0005  0.0185  0.0663
province2 0.0747  0.0123  6.0485 0.0000  0.0505  0.0989
province4 -0.0936  0.0137 -6.8339 0.0000 -0.1205 -0.0668
province5 -0.0545  0.0134 -4.0804 0.0000 -0.0806 -0.0283
province6 -0.0428  0.0146 -2.9299 0.0034 -0.0715 -0.0142
province7 -0.0487  0.0143 -3.4141 0.0006 -0.0767 -0.0208
urb_rur2  0.0307  0.0077  3.9854 0.0001  0.0156  0.0457

#-----
# Running probit regression
#-----
probit_model <- glm(hhsex ~ hh_size + urb_rur + province,
                  data = data, family = binomial(link = "probit"))
summary(probit_model)

```

③ Setting base/reference of factor variable.

④

coefficients are not directly interpretable.  
log of odd ratio

odd ratio:  $\frac{\text{Probability of an event occurring}}{\text{probability of an event not occurring}}$

⑤

⑥

```
Call:
glm(formula = hhsex ~ hh_size + urb_rur + province, family = binomial(link = "probit"),
    data = data)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-0.184775	0.038365	-4.816	1.46e-06 ***
hh_size	0.188169	0.007198	26.142	< 2e-16 ***
urb_rur2	0.097061	0.025233	3.847	0.000120 ***
province1	0.156245	0.042600	3.668	0.000245 ***
province2	0.267450	0.045381	5.893	3.78e-09 ***
province4	-0.291513	0.041813	-6.972	3.13e-12 ***
province5	-0.165594	0.041717	-3.969	7.20e-05 ***
province6	-0.125738	0.045751	-2.748	0.005991 **
province7	-0.151142	0.044465	-3.399	0.000676 ***

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 14830 on 12654 degrees of freedom  
Residual deviance: 13737 on 12646 degrees of freedom  
AIC: 13755

Number of Fisher Scoring iterations: 4

represents changes in z-score (not directly interpretable)

7

# Calculating marginal effects for probit regression

```
probit_margins <- margins(probit_model)
```

```
summary(probit_margins)
```

factor	AME	SE	z	p	lower	upper
hh_size	0.0579	0.0021	28.1014	0.0000	0.0538	0.0619
province1	0.0453	0.0122	3.7069	0.0002	0.0214	0.0693
province2	0.0746	0.0123	6.0510	0.0000	0.0505	0.0988
province4	-0.0959	0.0139	-6.8939	0.0000	-0.1231	-0.0686
province5	-0.0529	0.0134	-3.9466	0.0001	-0.0791	-0.0266
province6	-0.0397	0.0146	-2.7226	0.0065	-0.0683	-0.0111
province7	-0.0481	0.0143	-3.3679	0.0008	-0.0761	-0.0201
urb_rur2	0.0298	0.0077	3.8642	0.0001	0.0147	0.0448

## Task 8:

Using ~~MM1506~~ data (011-Affairs.RData), complete the following tasks.

- Load the **011-Affairs.RData**
- Tabulate the frequency of **affairs** variable from **Affairs** dataframe.
- Create a variable **ynaffairs** in **Affairs** dataframe such that the variable takes value 0 if no affairs and 1 if the person is involved in affairs.
- Set **ynaffairs** and **rating** variables as factor variables.
- Set '0' as reference for **ynaffairs** variable, '5' for **rating**, 'no' for **children**, and 'female' for **gender** variables.
- Fit a logistic regression model with **ynaffairs** as dependent variable and **gender**, **age**, **yearsmarried**, **children**, **rating** as independent variable.
- Calculate average marginal effect for each variables using the `margins()` function.

```
library(dplyr)
library(margins)
```

```
load('data/011-Affairs.RData')
```

```
table(Affairs$affairs)
 0    1    2    3    7   12
451  34   17   19   42   38
```

```
Affairs <- Affairs %>% mutate(ynaffair = case_when(affairs > 0 ~ 1, TRUE ~ 0),
```

```
      ynaffair = factor(ynaffair),
      rating = factor(rating))
```

```
table(Affairs$ynaffair)
 0    1
451 150
```

```
#setting 0 : No-Affairs as base/reference
```

```
Affairs$ynaffair <- relevel(Affairs$ynaffair, ref = '0')
```

```
#setting 5 : Very happy as base/reference
# 1 = very unhappy, 2 = somewhat unhappy, 3 = average, 4 = happier than
average, 5 = very happy.
```

✓ Affairs\$rating <- relevel(Affairs\$rating, ref = '5')

```
#setting no children as base/reference
```

✓ Affairs\$children <- relevel(Affairs\$children, ref = 'no')

```
#setting female as base/reference
```

✓ Affairs\$gender <- relevel(Affairs\$gender, ref = 'female')

✓ fit <- glm(ynaffair ~ gender  
+ age  
+ yearsmarried  
+ children  
+ rating,  
data=Affairs,  
family = binomial(link = "logit"))

```
summary(fit)
```

```
Call:
```

```
glm(formula = ynaffair ~ gender + age + yearsmarried + children +  
rating, family = binomial(link = "logit"), data = Affairs)
```

```
Coefficients:
```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.36331	0.45275	-3.011	0.00260 **
gendermale	0.38018	0.20644	1.842	0.06553 .
age	-0.04432	0.01789	-2.477	0.01324 *
yearsmarried	0.08127	0.03154	2.577	0.00997 **
childrenyes	0.32477	0.28716	1.131	0.25807
rating1	1.66252	0.55213	3.011	0.00260 **
rating2	1.64220	0.31872	5.152	2.57e-07 ***
rating3	0.76132	0.30044	2.534	0.01128 *
rating4	0.52336	0.25641	2.041	0.04124 *

```
---  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
(Dispersion parameter for binomial family taken to be 1)
```

```
Null deviance: 675.38 on 600 degrees of freedom  
Residual deviance: 622.26 on 592 degrees of freedom  
AIC: 640.26
```

```
Number of Fisher Scoring iterations: 4
```

✓ summary(margins(fit))

factor	AME	SE	z	p	lower	upper
age	-0.0075	0.0030	-2.5155	0.0119	-0.0134	-0.0017
childrenyes	0.0537	0.0459	1.1698	0.2421	-0.0363	0.1437
gendermale	0.0648	0.0350	1.8511	0.0642	-0.0038	0.1334
rating1	0.3274	0.1273	2.5714	0.0101	0.0778	0.5769
rating2	0.3225	0.0666	4.8427	0.0000	0.1920	0.4530
rating3	0.1248	0.0522	2.3916	0.0168	0.0225	0.2271
rating4	0.0803	0.0392	2.0468	0.0407	0.0034	0.1572
yearsmarried	0.0138	0.0053	2.6201	0.0088	0.0035	0.0242

## Session 11: Time-series analysis

### 11.1. Stationarity concept

- Stationarity refers to a time series whose statistical properties, such as mean, variance, and autocorrelation, remain constant over time.
- Non-stationary series are prone to spurious relationships.

### 11.2. Spurious relationships i.e. fake relationship

E040-spurious\_regression.R

```
library(haven)
library(dplyr)
library(tseries)
library(ggplot2)
```

① df <- read\_dta('data/012-pwt1001.dta')

```
#keeping real GDP of Nepal from 1960 onwards
npl <- df %>%
  filter(countrycode == 'NPL' & year >= 1960) %>%
  select(year, rgdpe) %>%
  rename(rgdpe_npl = rgdpe)
```

② Nepal and USA  
data sets making

```
#keeping real GDP of USA from 1960 onwards
usa <- df %>%
  filter(countrycode == 'USA' & year >= 1960) %>%
  select(year, rgdpe) %>%
  rename(rgdpe_usa = rgdpe)
```


```
#joining Nepal and USA data into one dataframe
df_npl_usa <- full_join(npl, usa, by = 'year')
```

③ Dataset merging

④ #Visual inspection of stationarity

```
ggplot() +
  geom_line(data = df_npl_usa, aes(x=year, y=rgdpe_npl), size = 1) +
  labs(title = 'Nepal GDP') +
  theme_bw()
```


Nepal GDP



⑤ ggplot() +

```
geom_line(data = df_npl_usa, aes(x=year, y=rgdpe_usa), size = 1) +
  labs(title = 'USA GDP') +
  theme_bw()
```

USA GDP



⑥ #Hypothesis testing of stationarity

```
adf.test(df_npl_usa$rgdpe_npl)
```

Augmented Dickey-Fuller Test

```
data: df_npl_usa$rgdpe_npl
Dickey-Fuller = 1.9811, Lag order = 3, p-value = 0.99
alternative hypothesis: stationary
```

```

6) adf.test(df_npl_usa$rgdpe_usa)
      Augmented Dickey-Fuller Test

data: df_npl_usa$rgdpe_usa
Dickey-Fuller = -1.1308, Lag order = 3, p-value = 0.9101
alternative hypothesis: stationary

```

```

7) #Running a regression (Spurious regression observed)
fit <- lm(formula = rgdpe_usa ~ rgdpe_npl, data = df_npl_usa)
summary(fit)
Call:
lm(formula = rgdpe_usa ~ rgdpe_npl, data = df_npl_usa)

Residuals:
    Min       1Q   Median       3Q      Max
-3982784 -1070155  -25168   853922  3683084

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 4384943.5   369375.8   11.87  <2e-16 ***
rgdpe_npl    230.4       10.7    21.54  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1741000 on 58 degrees of freedom
Multiple R-squared:  0.8888,    Adjusted R-squared:  0.8869
F-statistic: 463.8 on 1 and 58 DF, p-value: < 2.2e-16

```

```

7) # making series stationary by log differencing.
# Making series stationary and repeating the above steps
#-----
df_npl_usa <- df_npl_usa %>%
  mutate(dlrgdpe_npl = c(NA, diff(log(rgdpe_npl))),
         dlrgdpe_usa = c(NA, diff(log(rgdpe_usa)))) %>%
  na.omit()


```

Q: log difference  $\Rightarrow$  growth rate. How?

```

8) #Visual inspection of stationarity
ggplot() +
  geom_line(data = df_npl_usa, aes(x=year, y=dlrgdpe_npl), size = 1) +
  labs(title = 'Nepal GDP growth') +
  theme_bw()


```



```

ggplot() +
  geom_line(data = df_npl_usa, aes(x=year, y=dlrgdpe_usa), size = 1) +
  labs(title = 'USA GDP growth') +
  theme_bw()

```



⑨

#Hypothesis testing of stationarity

```
adf.test(df_npl_usa$dlrgdpe_npl)
Augmented Dickey-Fuller Test
```

```
data: df_npl_usa$dlrgdpe_npl
Dickey-Fuller = -3.236, Lag order = 3, p-value = 0.0904
alternative hypothesis: stationary
```

```
adf.test(df_npl_usa$dlrgdpe_usa)
Augmented Dickey-Fuller Test
```

```
data: df_npl_usa$dlrgdpe_usa
Dickey-Fuller = -4.4078, Lag order = 3, p-value = 0.01
alternative hypothesis: stationary
```

⑩

#Running a regression (no spurious regression observed)

```
fit <- lm(formula = dlrgdpe_usa ~ dlrgdpe_npl, data = df_npl_usa)
summary(fit)
```

```
Call:
lm(formula = dlrgdpe_usa ~ dlrgdpe_npl, data = df_npl_usa)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-0.052298 -0.006704  0.000871  0.014589  0.048773
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.031797   0.004109   7.738 1.87e-10 ***
dlrgdpe_npl -0.035680   0.070464  -0.506   0.615
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.02206 on 57 degrees of freedom
Multiple R-squared:  0.004478, Adjusted R-squared: -0.01299
F-statistic: 0.2564 on 1 and 57 DF, p-value: 0.6146
```

### 11.3. True relationships

#### E041-actual\_relationship.R

```
library(haven)
library(dplyr)
library(tseries)
library(ggplot2)
```

```
df <- read_dta('data/012-pwt1001.dta')
```

①


```
df <- filter(df, countrycode == 'NPL' & year >= 1960) %>% select(year,
rgdpe, ccon)
```

②


#Visual inspection of stationarity

```
ggplot() +
  geom_line(data = df, aes(x=year, y=rgdpe), size = 1) +
  labs(title = 'Nepal GDP') +
  theme_bw()
```





```
ggplot() +
  geom_line(data = df, aes(x=year, y=ccon), size = 1) +
  labs(title = 'Nepal Consumption (Private + Govt)') +
  theme_bw()
Nepal Consumption (Private + Govt)
```



⑧

#Hypothesis testing of stationarity

```
adf.test(df$rgdpe)
Augmented Dickey-Fuller Test
```

```
data: df$rgdpe
Dickey-Fuller = 1.9811, Lag order = 3, p-value = 0.99
alternative hypothesis: stationary
```

```
adf.test(df$ccon)
Augmented Dickey-Fuller Test
```

```
data: df$ccon
Dickey-Fuller = 0.87741, Lag order = 3, p-value = 0.99
alternative hypothesis: stationary
```

②

#Running a regression

```
fit <- lm(formula = rgdpe ~ ccon ,data = df)
summary(fit)
```



```
Call:
lm(formula = rgdpe ~ ccon, data = df)

Residuals:
    Min       1Q   Median       3Q      Max
-7032.0  -695.9  -295.4  1007.2  3166.8

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -157.77337   331.29087   -0.476   0.636
ccon         1.04900     0.01004  104.504 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1554 on 58 degrees of freedom
Multiple R-squared:  0.9947,    Adjusted R-squared:  0.9946
F-statistic: 1.092e+04 on 1 and 58 DF,  p-value: < 2.2e-16
```

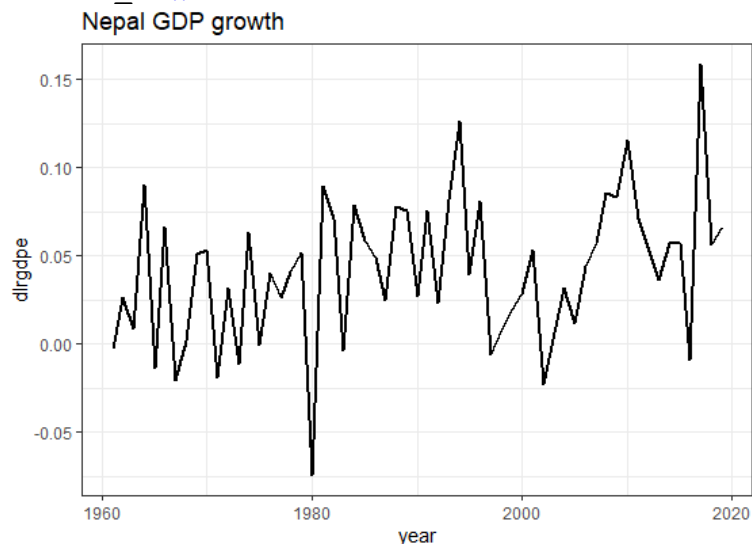
```
#-----
# Making series stationary and repeating the above steps
#-----
```

5


```
df <- df %>%
  mutate(dlrgdpe = c(NA, diff(log(rgdpe))),
         dlcccon = c(NA, diff(log(cccon)))) %>%
  na.omit()
```

6

```
#Visual inspection of stationarity
ggplot() +
  geom_line(data = df, aes(x=year, y=dlrgdpe), size = 1) +
  labs(title = 'Nepal GDP growth') +
  theme_bw()
```



```
ggplot() +
  geom_line(data = df, aes(x=year, y=dlcccon), size = 1) +
  labs(title = 'Nepal Consumption (Private + Govt) growth') +
  theme_bw()
```



```

⑦ #Hypothesis testing of stationarity
adf.test(df_npl_usa$dlrgdpe_npl)
      Augmented Dickey-Fuller Test

data: df_npl_usa$dlrgdpe_npl
Dickey-Fuller = -3.236, Lag order = 3, p-value = 0.0904
alternative hypothesis: stationary

adf.test(df_npl_usa$dlrgdpe_usa)
      Augmented Dickey-Fuller Test

data: df_npl_usa$dlrgdpe_usa
Dickey-Fuller = -4.4078, Lag order = 3, p-value = 0.01
alternative hypothesis: stationary

⑧ #Running a regression
fit <- lm(formula = dlrgdpe ~ dlcon ,data = df)
summary(fit)
Call:
lm(formula = dlrgdpe ~ dlcon, data = df)

Residuals:
    Min       1Q   Median       3Q      Max
-0.065718 -0.013749  0.000986  0.015554  0.101673

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.006254   0.004832   1.294    0.201
dlcon       0.880921   0.088548   9.949 4.54e-14 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.02507 on 57 degrees of freedom
Multiple R-squared:  0.6346,    Adjusted R-squared:  0.6281
F-statistic: 98.97 on 1 and 57 DF,  p-value: 4.542e-14

```

## Session 12: Stargazer for reporting regression results and the project work

### 12.1. stargazer

```

E042-stargazer.R
mtcars <- datasets::mtcars

④ model1 <- lm(mpg ~ hp, data = mtcars)
⑤ model2 <- lm(mpg ~ hp + drat, data = mtcars)
model3 <- lm(mpg ~ hp + drat + cyl + wt, data = mtcars)
model4 <- lm(hp ~ disp + carb, data = mtcars)

① library(stargazer)

#descriptive statistics table
⑤ stargazer(mtcars, type = 'text')
=====
Statistic N    Mean    St. Dev.  Min    Max
-----
mpg        32  20.091    6.027   10.400  33.900
cyl        32   6.188    1.786     4      8
disp       32 230.722  123.939   71.100 472.000
hp         32 146.688   68.563    52    335
drat       32   3.597    0.535    2.760   4.930
wt         32   3.217    0.978    1.513   5.424
qsec       32  17.849    1.787   14.500  22.900
vs         32   0.438    0.504     0      1
am         32   0.406    0.499     0      1
gear       32   3.688    0.738     3      5
carb       32   2.812    1.615     1      8
-----

⑤ #displaying regression models results in a single table
stargazer(model1, model2, model3, model4, type = "text")

```

Dependent variable:				
	(1)	mpg (2)	(3)	hp (4)
hp	-0.068*** (0.010)	-0.052*** (0.009)	-0.021 (0.013)	
drat		4.698*** (1.192)	0.818 (1.387)	
cyl			-0.762 (0.635)	
wt			-2.973*** (0.818)	
disp				0.324*** (0.043)
carb				21.999*** (3.298)
Constant	30.099*** (1.634)	10.790** (5.078)	34.496*** (7.441)	9.988 (11.614)
Observations	32	32	32	32
R2	0.602	0.741	0.845	0.852
Adjusted R2	0.589	0.723	0.822	0.842
Residual Std. Error	3.863 (df = 30)	3.170 (df = 29)	2.541 (df = 27)	27.244 (df = 29)
F Statistic	45.460*** (df = 1; 30)	41.522*** (df = 2; 29)	36.839*** (df = 4; 27)	83.665*** (df = 2; 29)
Note: *p<0.1; **p<0.05; ***p<0.01				

```

⑥ #defining the covariate and variable labels
stargazer(model1, model2, model3, model4, type = "text",
  digits = 2,
  covariate.labels = c('Gross horsepower (hp)',
    'Rear axle ratio (drat)',
    'Number of cylinders (cyl)',
    'Weight (1000 lbs) (wt)',
    'Displacement (cu.in.) (disp)',
    'Number of carburetors (carb)'),
  dep.var.labels = c("Miles/(US) gallon (mpg)", "Gross horsepower
    (hp)"),
  notes = "Standard errors are in parentheses.")

#export and save the result as html
⑦ stargazer(model1, model2, model3, model4, type = "html", out =
  'model_results.html',
  digits = 2,
  covariate.labels = c('Gross horsepower (hp)',
    'Rear axle ratio (drat)',
    'Number of cylinders (cyl)',
    'Weight (1000 lbs) (wt)',
    'Displacement (cu.in.) (disp)',
    'Number of carburetors (carb)'),
  dep.var.labels = c("Miles/(US) gallon (mpg)", "Gross horsepower
    (hp)"),
  notes = "Standard errors are in parentheses.")

```

	<i>Dependent variable:</i>			
	Miles/(US) gallon (mpg)			Gross horsepower (hp)
	(1)	(2)	(3)	(4)
Gross horsepower (hp)	-0.07*** (0.01)	-0.05*** (0.01)	-0.02 (0.01)	
Rear axle ratio (dart)		4.70*** (1.19)	0.82 (1.39)	
Number of cylinders (cyl)			-0.76 (0.64)	
Weight (1000 lbs) (wt)			-2.97*** (0.82)	
Displacement (cu.in.) (disp)				0.32*** (0.04)
Number of carburetors (carb)				22.00*** (3.30)
Constant	30.10*** (1.63)	10.79** (5.08)	34.50*** (7.44)	9.99 (11.61)
Observations	32	32	32	32
R <sup>2</sup>	0.60	0.74	0.85	0.85
Adjusted R <sup>2</sup>	0.59	0.72	0.82	0.84
Residual Std. Error	3.86 (df = 30)	3.17 (df = 29)	2.54 (df = 27)	27.24 (df = 29)
F Statistic	45.46*** (df = 1; 30)	41.52*** (df = 2; 29)	36.84*** (df = 4; 27)	83.66*** (df = 2; 29)
<i>Note:</i>			* p<0.1; ** p<0.05; *** p<0.01	
			Standard errors are in parentheses.	