

## Day 3 : Statistical analysis using R

### Session 9: Regression

#### 9.1. Fitting regression models with lm()

Symbols commonly used in R formulas

Symbol	Usage
~	Separates response variables on the left from the explanatory variables on the right. For example, a prediction of $y$ from $x$ , $z$ , and $w$ would be coded $y \sim x + z + w$ .
+	Separates predictor variables
:	Denotes an interaction between predictor variables. A prediction of $y$ from $x$ , $z$ , and the interaction between $x$ and $z$ would be coded $y \sim x + z + x:z$ .
*	A shortcut for denoting all possible interactions. The code $y \sim x * z * w$ expands to $y \sim x + z + w + x:z + x:w + z:w + x:z:w$ .
^	Denotes interactions up to a specified degree. The code $y \sim (x + z + w)^2$ expands to $y \sim x + z + w + x:z + x:w + z:w$ .
.	A placeholder for all other variables in the data frame except the dependent variable. For example, if a data frame contained the variables $x$ , $y$ , $z$ , and $w$ , then the code $y \sim .$ would expand to $y \sim x + z + w$ .
-	A minus sign removes a variable from the equation. For example, $y \sim (x + z + w)^2 - x:w$ expands to $y \sim x + z + w + x:z + z:w$ .
-1	Suppresses the intercept. For example, the formula $y \sim x - 1$ fits a regression of $y$ on $x$ and forces the line through the origin at $x=0$ .
I()	Elements within the parentheses are interpreted arithmetically. For example, $y \sim x + (z + w)^2$ expands to $y \sim x + z + w + z:w$ . In contrast, the code $y \sim x + I((z + w)^2)$ expands to $y \sim x + h$ , where $h$ is a new variable created by squaring the sum of $z$ and $w$ .
function	Mathematical functions can be used in formulas. For example, $\log(y) \sim x + z + w$ predicts $\log(y)$ from $x$ , $z$ , and $w$ .

#### E034-simple\_regression.R

```
# Set seed for reproducibility
set.seed(12345)

# Generate 1000 observations
n <- 1000

# Generate study_hours as uniform random numbers between 0 and 10
study_hours <- round(runif(n, min = 0, max = 10))

# Generate score as a linear function of study_hours with noise
score <- 50 + 5 * study_hours + rnorm(n, mean = 0, sd = 5)

# Combine into a data frame
df <- data.frame(study_hours, score)

# Perform linear regression
model <- lm(score ~ study_hours, data = df)

# Summarize the regression results
summary(model)
```


```
Call:
lm(formula = score ~ study_hours, data = df)

Residuals:
    Min       1Q   Median       3Q      Max
-16.612  -3.334  -0.018   3.509  16.737

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 49.94446    0.32742  152.54  <2e-16 ***
study_hours  4.99446    0.05571   89.64  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.035 on 998 degrees of freedom
Multiple R-squared:  0.8895,    Adjusted R-squared:  0.8894
F-statistic: 8036 on 1 and 998 DF, p-value: < 2.2e-16
```

```
# visualizing simple regression
library(ggplot2)
ggplot(data = df, aes(x=study_hours, y=score)) +
  geom_point() +
  geom_smooth(method = 'lm') +
  theme_bw()
```



## 9.2. Multiple regression

### E035-multiple\_regression.R

```
# Set seed for reproducibility
set.seed(12345)

# Generate 200 observations
n <- 200

# Generate age variable (cycles from 18 to 69)
age <- (1:n %% 52) + 18

# Generate educ_year variable (cycles from 0 to 17)
educ_year <- (1:n %% 18)

# Generate income variable with a linear relationship to age and educ_year,
plus noise
income <- 20000 + 800 * age + 3000 * educ_year + rnorm(n, mean = 0, sd =
2000)
```

```
# Combine into a data frame
df <- data.frame(age, educ_year, income)

# Regression with omitted variable
model_omitted <- lm(income ~ age, data = df)
summary(model_omitted)


Call:
lm(formula = income ~ age, data = df)

Residuals:
    Min       1Q   Median       3Q      Max
-29181.6 -13567.7   363.5  14563.5  27208.4

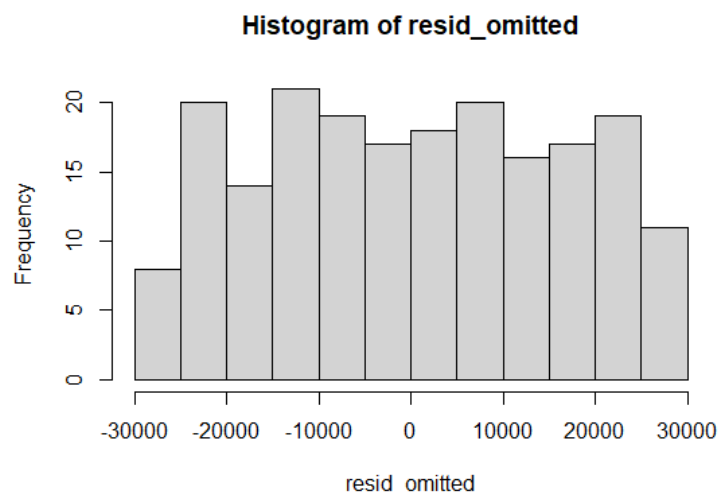
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 44039.53   3535.95   12.46  <2e-16 ***
age           835.97     78.13    10.70  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 16120 on 198 degrees of freedom
Multiple R-squared:  0.3663,    Adjusted R-squared:  0.3631
F-statistic: 114.5 on 1 and 198 DF,  p-value: < 2.2e-16

# Residual diagnostics for omitted variable model
plot(model_omitted, which = 1) # Residual vs Fitted plot
```



```
resid_omitted <- residuals(model_omitted)
hist(resid_omitted)
```




```
shapiro.test(resid_omitted) # Shapiro-Wilk test for normality [H0: normally distributed]
```

### Shapiro-Wilk normality test


```
data: resid_omitted  
W = 0.95517, p-value = 6.14e-06
```

```
#-----  
# Multiple regression with correct specification  
model_correct <- lm(income ~ age + educ_year, data = df)  
summary(model_correct)
```


**Histogram of resid\_omitted**



```
# Residual diagnostics for correctly specified model  
plot(model_correct, which = 1) # Residual vs Fitted plot
```



```
lm(income ~ age + educ_year)  
resid_correct <- residuals(model_correct)  
hist(resid_correct)
```



### 9.3. Polynomial regression

```

E036-polynomial_regression.R
library(ggplot2)

mtcars <- datasets::mtcars

#-----
#simple regression
#-----
fit <- lm(data = mtcars, formula = mpg ~ hp) # mpg: Miles/(US) gallon, hp:
Gross horsepower
summary(fit) #R-squared : 0.6024, Residual standard error: 3.863
Call:
lm(formula = mpg ~ hp, data = mtcars)


Residuals:
    Min       1Q   Median       3Q      Max
-5.7121 -2.1122 -0.8854  1.5819  8.2360

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 30.09886   1.63392  18.421 < 2e-16 ***
hp          -0.06823   0.01012  -6.742 1.79e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.863 on 30 degrees of freedom
Multiple R-squared:  0.6024,    Adjusted R-squared:  0.5892
F-statistic: 45.46 on 1 and 30 DF,  p-value: 1.788e-07

ggplot(mtcars, aes(x = hp, y = mpg)) +
  geom_point() +
  stat_smooth(method = 'lm', formula = y ~ x, color = 'red', se = FALSE) +
  theme_bw()

```




```
#-----
#Polynomial regression
#-----
fit <- lm(data = mtcars, formula = mpg ~ hp + I(hp^2))
summary(fit) #R-squared : 0.7561, Residual standard error: 3.077
Call:
lm(formula = mpg ~ hp + I(hp^2), data = mtcars)

Residuals:
    Min       1Q   Median       3Q      Max
-4.5512 -1.6027 -0.6977  1.5509  8.7213

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  4.041e+01  2.741e+00  14.744 5.23e-15 ***
hp          -2.133e-01  3.488e-02  -6.115 1.16e-06 ***
I(hp^2)       4.208e-04  9.844e-05   4.275 0.000189 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.077 on 29 degrees of freedom
Multiple R-squared:  0.7561,    Adjusted R-squared:  0.7393
F-statistic: 44.95 on 2 and 29 DF,  p-value: 1.301e-09
```

```
ggplot(mtcars, aes(x = hp, y = mpg)) +
  geom_point() +
  stat_smooth(method = 'lm', formula = y ~ x + I(x^2), color = 'red', se =
FALSE) +
  theme_bw()
```



## 9.4. Regression with interaction term

### E037-regression\_with\_interaction.R

```
mtcars <- datasets::mtcars

#generating a new interaction term hp * wt
mtcars$hp_wt <- mtcars$hp * mtcars$wt

fit <- lm(mpg ~ hp + wt + hp_wt, data=mtcars)
summary(fit)

Call:
lm(formula = mpg ~ hp + wt + hp_wt, data = mtcars)

Residuals:
    Min       1Q   Median       3Q      Max
-3.0632 -1.6491 -0.7362  1.4211  4.5513

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  49.80842     3.60516   13.816 5.01e-14 ***
hp          -0.12010     0.02470    -4.863 4.04e-05 ***
wt          -8.21662     1.26971    -6.471 5.20e-07 ***
hp_wt         0.02785     0.00742     3.753 0.000811 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.153 on 28 degrees of freedom
Multiple R-squared:  0.8848,    Adjusted R-squared:  0.8724
F-statistic: 71.66 on 3 and 28 DF,  p-value: 2.981e-13

#OR

fit <- lm(mpg ~ hp + wt + hp:wt, data=mtcars)
summary(fit)

Call:
lm(formula = mpg ~ hp + wt + hp:wt, data = mtcars)

Residuals:
    Min       1Q   Median       3Q      Max
-3.0632 -1.6491 -0.7362  1.4211  4.5513

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  49.80842     3.60516   13.816 5.01e-14 ***
hp          -0.12010     0.02470    -4.863 4.04e-05 ***
wt          -8.21662     1.26971    -6.471 5.20e-07 ***
hp:wt         0.02785     0.00742     3.753 0.000811 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.153 on 28 degrees of freedom
Multiple R-squared:  0.8848,    Adjusted R-squared:  0.8724
F-statistic: 71.66 on 3 and 28 DF,  p-value: 2.981e-13

##* -----
##* A significant coefficient of interaction term indicates that
##* the relationship between mpg and hp varies by wt. Similarly,
##* the relationship between mpg and wt varies by hp.
##* -----

# d(mpg)/d(hp) = - 0.12010 + 0.02785 * wt
wt = 1
print(- 0.12010 + 0.02785 * wt) #-0.09225

wt = 2
print(- 0.12010 + 0.02785 * wt) #-0.0644

wt = 3
print(- 0.12010 + 0.02785 * wt) #-0.03655

# d(mpg)/d(wt) = - 8.21662 + 0.02785 * hp
hp = 100
```

```
print(- 8.21662 + 0.02785 * hp) #-5.43162

hp = 150
print(- 8.21662 + 0.02785 * hp) #-4.03912

hp = 200
print(- 8.21662 + 0.02785 * hp) #-2.64662
```

## 9.5. Logarithmic regression

Model	Equation	Interpretation of $\beta_1$
Log-Log	$\log(y) = \beta_0 + \beta_1 \log(x)$	Elasticity: 1% change in $x$ leads to $\beta_1\%$ change in $y$ .
Log-Linear	$\log(y) = \beta_0 + \beta_1 x$	Semi-elasticity: 1-unit change in $x$ leads to $(\exp(\beta_1) - 1) \times 100\%$ change in $y$ .
Linear-Log	$y = \beta_0 + \beta_1 \log(x)$	1% change in $x$ leads to $\beta_1/100$ unit change in $y$ .
Linear-Linear	$y = \beta_0 + \beta_1 x$	1-unit change in $x$ leads to $\beta_1$ unit change in $y$ .

### E038-logarithmic\_regression.R

```
#-----
# Log-Log Regression
#-----
# Load data
mtcars <- datasets::mtcars

# Log-log regression
model_loglog <- lm(log(mpg) ~ log(displacement), data = mtcars)
summary(model_loglog)
Call:
lm(formula = log(mpg) ~ log(displacement), data = mtcars)

Residuals:
    Min       1Q   Median       3Q      Max
-0.22758 -0.08874 -0.00791  0.07970  0.32143

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  5.38097    0.20803   25.87 < 2e-16 ***
log(displacement) -0.45857    0.03913  -11.72 1.01e-12 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1282 on 30 degrees of freedom
Multiple R-squared:  0.8207,    Adjusted R-squared:  0.8148
F-statistic: 137.3 on 1 and 30 DF,  p-value: 1.006e-12

# A 1% increase in Displacement (cu.in.) reduces Miles/(US) gallon by
~0.46%.

#-----
# Log-Linear Regression
#-----
# Log-linear regression
model_loglin <- lm(log(mpg) ~ hp, data = mtcars)
summary(model_loglin)
```



```

Call:
lm(formula = log(mpg) ~ hp, data = mtcars)

Residuals:
    Min       1Q   Median       3Q      Max
-0.41577 -0.06583 -0.01737  0.09827  0.39621

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  3.4604669   0.0785838   44.035 < 2e-16 ***
hp          -0.0034287   0.0004867   -7.045 7.85e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1858 on 30 degrees of freedom
Multiple R-squared:  0.6233,    Adjusted R-squared:  0.6107
F-statistic: 49.63 on 1 and 30 DF,  p-value: 7.853e-08

# A 1-unit increase in horsepower reduces MPG by ~0.34% (exp(-0.0034287) - 1
# ≈ -0.003422829).

#-----
# Linear-Log Regression
#-----

# Load data
trees <- datasets::trees

# Linear-log regression
model_linlog <- lm(Volume ~ log(Girth), data = trees)
summary(model_linlog)
Call:
lm(formula = Volume ~ log(Girth), data = trees)

Residuals:
    Min       1Q   Median       3Q      Max
-9.7246 -3.5312 -0.9174  3.2154 15.8780

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -138.973     11.439  -12.15 6.71e-13 ***
log(Girth)   66.141       4.455   14.85 4.38e-15 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.701 on 29 degrees of freedom
Multiple R-squared:  0.8837,    Adjusted R-squared:  0.8797
F-statistic: 220.4 on 1 and 29 DF,  p-value: 4.381e-15

# A 1% increase in girth increases volume by ~ 0.66 units (66.141 / 100).

```

## Session 10: Logistic regression

### 11.1. Logistic regression

Logistic regression is useful when you're predicting a binary outcome from a set of continuous and/or categorical predictor variables.

#### E039-logistic\_regression.R

```

# Load necessary libraries
library(haven) # For reading SPSS files
library(dplyr) # For data manipulation
library(margins) # For calculating marginal effects

# Import SPSS file from the URL
data <- read_spss('data/010-hh.sav')

# Dropping missing values in HHSEX
data <- data %>% filter(!is.na(HHSEX))

# Creating new variables
data <- data %>%
  mutate(
    hh_size = HH48, # HH member size variable
    urb_rur = factor(HH6), # 1=Urban 2=Rural
    province = factor(HH7), # Province number

```

```

    hhsex = factor(HHSEX) # 1=Male 2=Female
  )

#setting 1=Urban as reference/base
data$urb_rur <- relevel(data$urb_rur, ref = '1')

#setting 2=Female as reference/base
data$hhsex <- relevel(data$hhsex, ref = '2')

#setting province 3 as base category/reference level
data$province <- relevel(data$province, ref = '3')

#-----
# Running logistic regression
#-----
logit_model <- glm(hhsex ~ hh_size + urb_rur + province,
                  data = data, family = binomial(link = "logit"))
summary(logit_model)
Call:
glm(formula = hhsex ~ hh_size + urb_rur + province, family = binomial(link = "logit"),
    data = data)

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -0.39332    0.06488  -6.062 1.35e-09 ***
hh_size      0.33308    0.01274  26.145 < 2e-16 ***
urb_rur2     0.16929    0.04273   3.962 7.44e-05 ***
province1    0.24989    0.07272   3.436 0.000590 ***
province2    0.46344    0.07949   5.830 5.53e-09 ***
province4   -0.47811    0.06912  -6.917 4.61e-12 ***
province5   -0.28721    0.06988  -4.110 3.95e-05 ***
province6   -0.22815    0.07692  -2.966 0.003017 **
province7   -0.25820    0.07476  -3.454 0.000553 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 14830  on 12654  degrees of freedom
Residual deviance: 13722  on 12646  degrees of freedom
AIC: 13740

Number of Fisher Scoring iterations: 4

# Calculating marginal effects for logistic regression
logit_margins <- margins(logit_model)
summary(logit_margins)
  factor    AME    SE      z      p   lower   upper
hh_size  0.0606 0.0021 28.5476 0.0000  0.0564  0.0647
province1 0.0424 0.0122  3.4806 0.0005  0.0185  0.0663
province2 0.0747 0.0123  6.0485 0.0000  0.0505  0.0989
province4 -0.0936 0.0137 -6.8339 0.0000 -0.1205 -0.0668
province5 -0.0545 0.0134 -4.0804 0.0000 -0.0806 -0.0283
province6 -0.0428 0.0146 -2.9299 0.0034 -0.0715 -0.0142
province7 -0.0487 0.0143 -3.4141 0.0006 -0.0767 -0.0208
urb_rur2  0.0307 0.0077  3.9854 0.0001  0.0156  0.0457

#-----
# Running probit regression
#-----
probit_model <- glm(hhsex ~ hh_size + urb_rur + province,
                  data = data, family = binomial(link = "probit"))
summary(probit_model)

```

```
Call:
glm(formula = hhsex ~ hh_size + urb_rur + province, family = binomial(link = "probit"),
    data = data)

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -0.184775   0.038365  -4.816 1.46e-06 ***
hh_size      0.188169   0.007198  26.142 < 2e-16 ***
urb_rur2     0.097061   0.025233   3.847 0.000120 ***
province1    0.156245   0.042600   3.668 0.000245 ***
province2    0.267450   0.045381   5.893 3.78e-09 ***
province4   -0.291513   0.041813  -6.972 3.13e-12 ***
province5   -0.165594   0.041717  -3.969 7.20e-05 ***
province6   -0.125738   0.045751  -2.748 0.005991 **
province7   -0.151142   0.044465  -3.399 0.000676 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 14830  on 12654  degrees of freedom
Residual deviance: 13737  on 12646  degrees of freedom
AIC: 13755

Number of Fisher Scoring iterations: 4

# Calculating marginal effects for probit regression
probit_margins <- margins(probit_model)
summary(probit_margins)
  factor    AME    SE      z      p   lower   upper
hh_size  0.0579 0.0021 28.1014 0.0000  0.0538  0.0619
province1 0.0453 0.0122  3.7069 0.0002  0.0214  0.0693
province2 0.0746 0.0123  6.0510 0.0000  0.0505  0.0988
province4 -0.0959 0.0139 -6.8939 0.0000 -0.1231 -0.0686
province5 -0.0529 0.0134 -3.9466 0.0001 -0.0791 -0.0266
province6 -0.0397 0.0146 -2.7226 0.0065 -0.0683 -0.0111
province7 -0.0481 0.0143 -3.3679 0.0008 -0.0761 -0.0201
urb_rur2  0.0298 0.0077  3.8642 0.0001  0.0147  0.0448
```

## Task 8:

Using NMICS6 data (011-Affairs.RData), complete the following tasks.

- Load the **011-Affairs.RData**
- Tabulate the frequency of **affairs** variable from **Affairs** dataframe.
- Create a variable **ynaffairs** in **Affairs** dataframe such that the variable takes value 0 if no affairs and 1 if the person is involved in affairs.
- Set **ynaffairs** and **rating** variables as factor variables.
- Set '0' as reference for **ynaffairs** variable, '5' for **rating**, 'no' for **children**, and 'female' for **gender** variables.
- Fit a logistic regression model with **ynaffairs** as dependent variable and **gender**, **age**, **yearsmarried**, **children**, **rating** as independent variable.
- Calculate average marginal effect for each variables using the `margins()` function.

```
library(dplyr)
library(margins)

load('data/011-Affairs.RData')
table(Affairs$affairs)
  0    1    2    3    7   12
451  34   17   19   42   38

Affairs <- Affairs %>% mutate(ynaffair = case_when(affairs > 0 ~ 1, TRUE ~
0),
                             ynaffair = factor(ynaffair),
                             rating = factor(rating))

table(Affairs$ynaffair)
  0    1
451 150

#setting 0 : No-Affairs as base/reference
Affairs$ynaffair <- relevel(Affairs$ynaffair, ref = '0')
```

```

#setting 5 : Very happy as base/reference
# 1 = very unhappy, 2 = somewhat unhappy, 3 = average, 4 = happier than
average, 5 = very happy.
Affairs$rating <- relevel(Affairs$rating, ref = '5')

#setting no children as base/reference
Affairs$children <- relevel(Affairs$children, ref = 'no')

#setting female as base/reference
Affairs$gender <- relevel(Affairs$gender, ref = 'female')

fit <- glm(ynaffair ~ gender
          + age
          + yearsmarried
          + children
          + rating,
          data=Affairs,
          family = binomial(link = "logit"))

summary(fit)
Call:
glm(formula = ynaffair ~ gender + age + yearsmarried + children +
    rating, family = binomial(link = "logit"), data = Affairs)

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -1.36331    0.45275  -3.011  0.00260 **
gendermale   0.38018    0.20644   1.842  0.06553 .
age         -0.04432    0.01789  -2.477  0.01324 *
yearsmarried 0.08127    0.03154   2.577  0.00997 **
childrenyes  0.32477    0.28716   1.131  0.25807
rating1      1.66252    0.55213   3.011  0.00260 **
rating2      1.64220    0.31872   5.152 2.57e-07 ***
rating3      0.76132    0.30044   2.534  0.01128 *
rating4      0.52336    0.25641   2.041  0.04124 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 675.38  on 600  degrees of freedom
Residual deviance: 622.26  on 592  degrees of freedom
AIC: 640.26

Number of Fisher Scoring iterations: 4

summary(margins(fit))
      factor      AME      SE      z      p      lower      upper
age -0.0075 0.0030 -2.5155 0.0119 -0.0134 -0.0017
childrenyes 0.0537 0.0459 1.1698 0.2421 -0.0363 0.1437
gendermale  0.0648 0.0350 1.8511 0.0642 -0.0038 0.1334
rating1     0.3274 0.1273 2.5714 0.0101 0.0778 0.5769
rating2     0.3225 0.0666 4.8427 0.0000 0.1920 0.4530
rating3     0.1248 0.0522 2.3916 0.0168 0.0225 0.2271
rating4     0.0803 0.0392 2.0468 0.0407 0.0034 0.1572
yearsmarried 0.0138 0.0053 2.6201 0.0088 0.0035 0.0242

```

## Session 11: Time-series analysis

### 11.1. Stationarity concept

- Stationarity refers to a time series whose statistical properties, such as mean, variance, and autocorrelation, remain constant over time.
- Non-stationary series are prone to spurious relationships.

### 11.2. Spurious relationships

#### E040-spurious\_regression.R

```

library(haven)
library(dplyr)
library(tseries)
library(ggplot2)

df <- read_dta('data/012-pwt1001.dta')

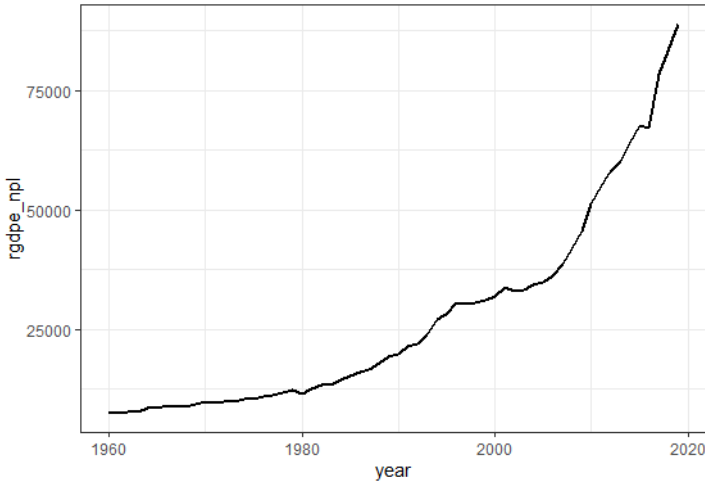
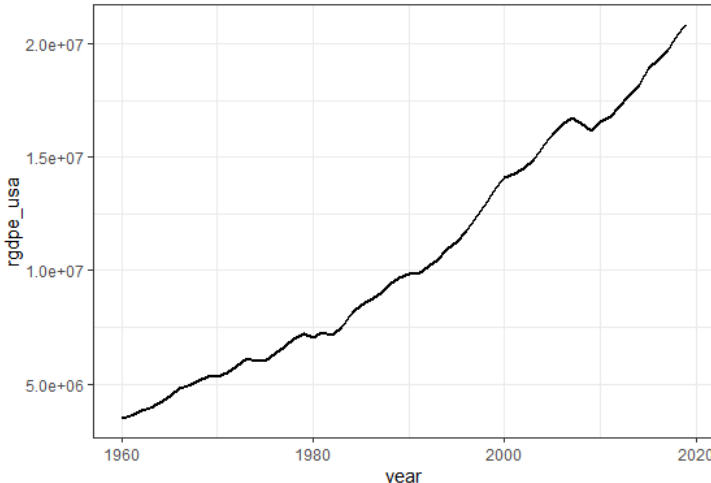
```

```

#keeping real GDP of Nepal from 1960 onwards
npl <- df %>%
  filter(countrycode == 'NPL' & year >= 1960) %>%
  select(year, rgdpe) %>%
  rename(rgdpe_npl = rgdpe)

#keeping real GDP of USA from 1960 onwards
usa <- df %>%
  filter(countrycode == 'USA' & year >= 1960) %>%
  select(year, rgdpe) %>%
  rename(rgdpe_usa = rgdpe)

#joining Nepal and USA data into one dataframe
df_npl_usa <- full_join(npl, usa, by = 'year')

#Visual inspection of stationarity
ggplot() +
  geom_line(data = df_npl_usa, aes(x=year, y=rgdpe_npl), size = 1) +
  labs(title = 'Nepal GDP') +
  theme_bw()
  Nepal GDP

  ggplot() +
  geom_line(data = df_npl_usa, aes(x=year, y=rgdpe_usa), size = 1) +
  labs(title = 'USA GDP') +
  theme_bw()
  USA GDP

  #Hypothesis testing of stationarity
  adf.test(df_npl_usa$rgdpe_npl)
  Augmented Dickey-Fuller Test

  data: df_npl_usa$rgdpe_npl
  Dickey-Fuller = 1.9811, Lag order = 3, p-value = 0.99
  alternative hypothesis: stationary

```

```

adf.test(df_npl_usa$rgdpe_usa)
      Augmented Dickey-Fuller Test

data: df_npl_usa$rgdpe_usa
Dickey-Fuller = -1.1308, Lag order = 3, p-value = 0.9101
alternative hypothesis: stationary

#Running a regression (Spurious regression observed)
fit <- lm(formula = rgdpe_usa ~ rgdpe_npl ,data = df_npl_usa)
summary(fit)
Call:
lm(formula = rgdpe_usa ~ rgdpe_npl, data = df_npl_usa)

Residuals:
    Min       1Q   Median       3Q      Max
-3982784 -1070155  -25168    853922   3683084

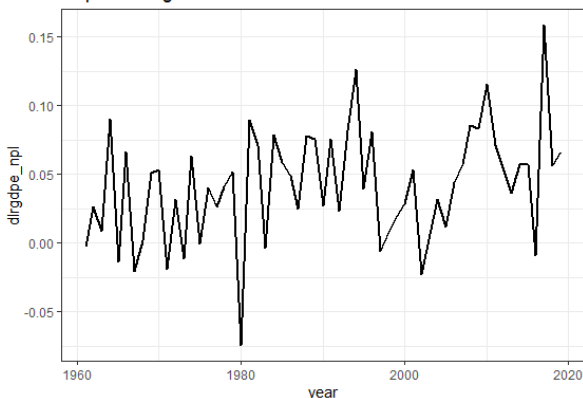
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 4384943.5   369375.8   11.87  <2e-16 ***
rgdpe_npl    230.4       10.7    21.54  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1741000 on 58 degrees of freedom
Multiple R-squared:  0.8888,    Adjusted R-squared:  0.8869
F-statistic: 463.8 on 1 and 58 DF,  p-value: < 2.2e-16

#-----
# Making series stationary and repeating the above steps
#-----
df_npl_usa <- df_npl_usa %>%
  mutate(dlrgdpe_npl = c(NA,diff(log(rgdpe_npl))),
         dlrgdpe_usa = c(NA,diff(log(rgdpe_usa)))) %>%
  na.omit()

#Visual inspection of stationarity
ggplot() +
  geom_line(data = df_npl_usa, aes(x=year, y=dlrgdpe_npl), size = 1) +
  labs(title = 'Nepal GDP growth') +
  theme_bw()


```



```

ggplot() +
  geom_line(data = df_npl_usa, aes(x=year, y=dlrgdpe_usa), size = 1) +
  labs(title = 'USA GDP growth') +
  theme_bw()

```



```
#Hypothesis testing of stationarity
adf.test(df_npl_usa$dlrgdpe_npl)
      Augmented Dickey-Fuller Test

data:  df_npl_usa$dlrgdpe_npl
Dickey-Fuller = -3.236, Lag order = 3, p-value = 0.0904
alternative hypothesis: stationary

adf.test(df_npl_usa$dlrgdpe_usa)
      Augmented Dickey-Fuller Test

data:  df_npl_usa$dlrgdpe_usa
Dickey-Fuller = -4.4078, Lag order = 3, p-value = 0.01
alternative hypothesis: stationary

#Running a regression (no spurious regression observed)
fit <- lm(formula = dlrgdpe_usa ~ dlrgdpe_npl ,data = df_npl_usa)
summary(fit)
Call:
lm(formula = dlrgdpe_usa ~ dlrgdpe_npl, data = df_npl_usa)

Residuals:
    Min       1Q   Median       3Q      Max
-0.052298 -0.006704  0.000871  0.014589  0.048773

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.031797   0.004109   7.738 1.87e-10 ***
dlrgdpe_npl -0.035680   0.070464  -0.506   0.615
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.02206 on 57 degrees of freedom
Multiple R-squared:  0.004478, Adjusted R-squared:  -0.01299
F-statistic: 0.2564 on 1 and 57 DF, p-value: 0.6146
```

### 11.3. True relationships


#### E041-actual\_relationship.R

```
library(haven)
library(dplyr)
library(tseries)
library(ggplot2)


df <- read_dta('data/012-pwt1001.dta')

df <- filter(df, countrycode == 'NPL' & year >= 1960) %>% select(year,
rgdpe, ccon)

#Visual inspection of stationarity
ggplot() +
  geom_line(data = df, aes(x=year, y=rgdpe), size = 1) +
  labs(title = 'Nepal GDP') +
  theme_bw()
```



```
ggplot() +
  geom_line(data = df, aes(x=year, y=ccon), size = 1) +
  labs(title = 'Nepal Consumption (Private + Govt)') +
  theme_bw()
Nepal Consumption (Private + Govt)
```



#Hypothesis testing of stationarity

```
adf.test(df$rgdpe)
Augmented Dickey-Fuller Test
```

```
data: df$rgdpe
Dickey-Fuller = 1.9811, Lag order = 3, p-value = 0.99
alternative hypothesis: stationary
```

```
adf.test(df$ccon)
Augmented Dickey-Fuller Test
```

```
data: df$ccon
Dickey-Fuller = 0.87741, Lag order = 3, p-value = 0.99
alternative hypothesis: stationary
```

#Running a regression

```
fit <- lm(formula = rgdpe ~ ccon ,data = df)
summary(fit)
```



```

Call:
lm(formula = rgdpe ~ ccon, data = df)

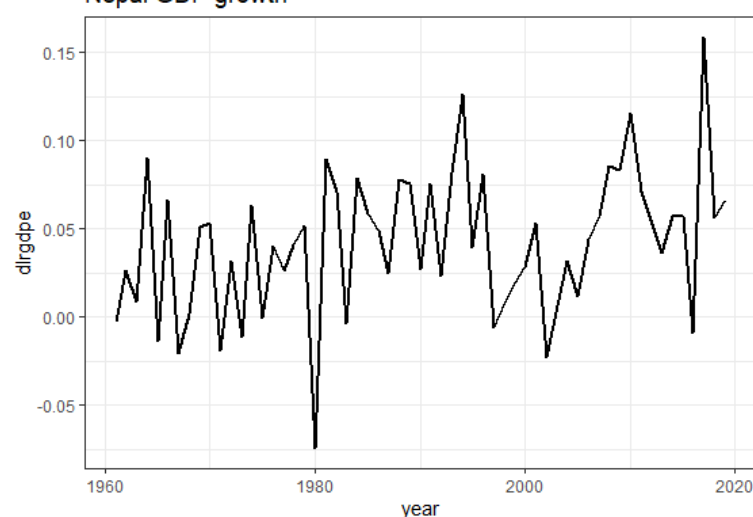
Residuals:
    Min       1Q   Median       3Q      Max
-7032.0  -695.9  -295.4  1007.2  3166.8

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -157.77337   331.29087   -0.476   0.636
ccon          1.04900     0.01004  104.504 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1554 on 58 degrees of freedom
Multiple R-squared:  0.9947,    Adjusted R-squared:  0.9946
F-statistic: 1.092e+04 on 1 and 58 DF,  p-value: < 2.2e-16
#-----
# Making series stationary and repeating the above steps
#-----
df <- df %>%
  mutate(dlrgdpe = c(NA,diff(log(rgdpe))),
         dlccon = c(NA,diff(log(ccon)))) %>%
  na.omit()

#Visual inspection of stationarity
ggplot() +
  geom_line(data = df, aes(x=year, y=dlrgdpe), size = 1) +
  labs(title = 'Nepal GDP growth') +
  theme_bw()

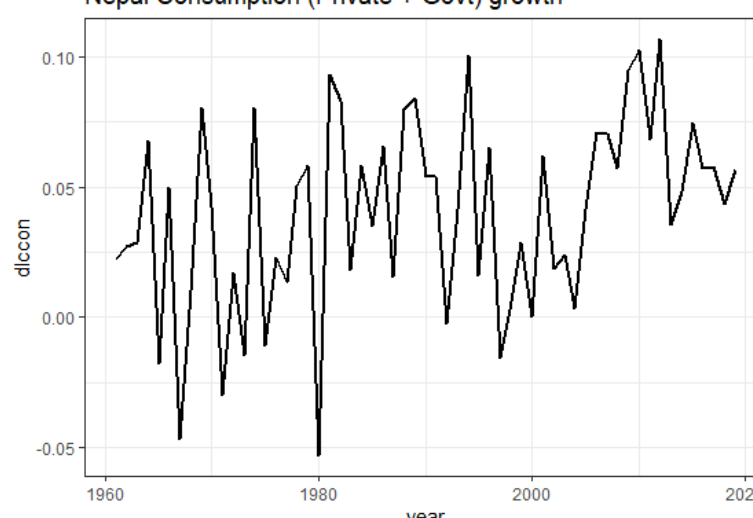
```



```

ggplot() +
  geom_line(data = df, aes(x=year, y=dlccon), size = 1) +
  labs(title = 'Nepal Consumption (Private + Govt) growth') +
  theme_bw()

```



```
#Hypothesis testing of stationarity
adf.test(df_npl_usa$dlrgdpe_npl)
      Augmented Dickey-Fuller Test

data: df_npl_usa$dlrgdpe_npl
Dickey-Fuller = -3.236, Lag order = 3, p-value = 0.0904
alternative hypothesis: stationary

adf.test(df_npl_usa$dlrgdpe_usa)
      Augmented Dickey-Fuller Test

data: df_npl_usa$dlrgdpe_usa
Dickey-Fuller = -4.4078, Lag order = 3, p-value = 0.01
alternative hypothesis: stationary

#Running a regression
fit <- lm(formula = dlrgdpe ~ dlcon ,data = df)
summary(fit)
Call:
lm(formula = dlrgdpe ~ dlcon, data = df)

Residuals:
    Min       1Q   Median       3Q      Max
-0.065718 -0.013749  0.000986  0.015554  0.101673

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.006254   0.004832   1.294   0.201
dlcon        0.880921   0.088548   9.949 4.54e-14 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.02507 on 57 degrees of freedom
Multiple R-squared:  0.6346,    Adjusted R-squared:  0.6281
F-statistic: 98.97 on 1 and 57 DF,  p-value: 4.542e-14
```

## Session 12: Stargazer for reporting regression results and the project work

### 12.1. stargazer

```
E042-stargazer.R
mtcars <- datasets::mtcars

model1 <- lm(mpg ~ hp, data = mtcars)
model2 <- lm(mpg ~ hp + drat, data = mtcars)
model3 <- lm(mpg ~ hp + drat + cyl + wt, data = mtcars)
model4 <- lm(hp ~ disp + carb, data = mtcars)

library(stargazer)

#descriptive statistics table
stargazer(mtcars, type = 'text')
=====
Statistic N    Mean    St. Dev.  Min    Max
-----
mpg       32  20.091    6.027   10.400  33.900
cyl       32   6.188    1.786     4      8
disp      32 230.722  123.939   71.100 472.000
hp        32 146.688   68.563    52    335
drat      32   3.597    0.535    2.760   4.930
wt        32   3.217    0.978    1.513   5.424
qsec      32  17.849    1.787   14.500  22.900
vs        32   0.438    0.504     0      1
am        32   0.406    0.499     0      1
gear      32   3.688    0.738     3      5
carb      32   2.812    1.615     1      8
-----

#displaying regression models results in a single table
stargazer(model1, model2, model3, model4, type = "text")
```

Dependent variable:				
	(1)	mpg (2)	(3)	hp (4)
hp	-0.068*** (0.010)	-0.052*** (0.009)	-0.021 (0.013)	
drat		4.698*** (1.192)	0.818 (1.387)	
cyl			-0.762 (0.635)	
wt			-2.973*** (0.818)	
disp				0.324*** (0.043)
carb				21.999*** (3.298)
Constant	30.099*** (1.634)	10.790** (5.078)	34.496*** (7.441)	9.988 (11.614)
Observations	32	32	32	32
R2	0.602	0.741	0.845	0.852
Adjusted R2	0.589	0.723	0.822	0.842
Residual Std. Error	3.863 (df = 30)	3.170 (df = 29)	2.541 (df = 27)	27.244 (df = 29)
F Statistic	45.460*** (df = 1; 30)	41.522*** (df = 2; 29)	36.839*** (df = 4; 27)	83.665*** (df = 2; 29)
Note: *p<0.1; **p<0.05; ***p<0.01				
<pre> #defining the covariate and variable labels stargazer(model1, model2, model3, model4, type = "text",   digits = 2,   covariate.labels = c('Gross horsepower (hp)',                         'Rear axle ratio (drat)',                         'Number of cylinders (cyl)',                         'Weight (1000 lbs) (wt)',                         'Displacement (cu.in.) (disp)',                         'Number of carburetors (carb)'),   dep.var.labels = c("Miles/(US) gallon (mpg)", "Gross horsepower (hp)"),   notes = "Standard errors are in parentheses.")  #export and save the result as html stargazer(model1, model2, model3, model4, type = "html", out = 'model_results.html',   digits = 2,   covariate.labels = c('Gross horsepower (hp)',                         'Rear axle ratio (drat)',                         'Number of cylinders (cyl)',                         'Weight (1000 lbs) (wt)',                         'Displacement (cu.in.) (disp)',                         'Number of carburetors (carb)'),   dep.var.labels = c("Miles/(US) gallon (mpg)", "Gross horsepower (hp)"),   notes = "Standard errors are in parentheses.") </pre>				

	<i>Dependent variable:</i>			
	Miles/(US) gallon (mpg)			Gross horsepower (hp)
	(1)	(2)	(3)	(4)
Gross horsepower (hp)	-0.07*** (0.01)	-0.05*** (0.01)	-0.02 (0.01)	
Rear axle ratio (dart)		4.70*** (1.19)	0.82 (1.39)	
Number of cylinders (cyl)			-0.76 (0.64)	
Weight (1000 lbs) (wt)			-2.97*** (0.82)	
Displacement (cu.in.) (disp)				0.32*** (0.04)
Number of carburetors (carb)				22.00*** (3.30)
Constant	30.10*** (1.63)	10.79** (5.08)	34.50*** (7.44)	9.99 (11.61)
Observations	32	32	32	32
R <sup>2</sup>	0.60	0.74	0.85	0.85
Adjusted R <sup>2</sup>	0.59	0.72	0.82	0.84
Residual Std. Error	3.86 (df = 30)	3.17 (df = 29)	2.54 (df = 27)	27.24 (df = 29)
F Statistic	45.46*** (df = 1; 30)	41.52*** (df = 2; 29)	36.84*** (df = 4; 27)	83.66*** (df = 2; 29)
<i>Note:</i>			* p<0.1; ** p<0.05; *** p<0.01	
			Standard errors are in parentheses.	